

# Advertising Data Analysis Using Theory of Constraints

## Environment Setup & Data Loading

```
In [45]: import os
import warnings

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

warnings.filterwarnings("ignore")
plt.style.use("default")
sns.set_context("notebook")
DATA_PATH = "advertising.csv"

if not os.path.exists(DATA_PATH):
    raise FileNotFoundError(f"Dataset not found at path: {DATA_PATH}")

df = pd.read_csv(DATA_PATH)
df.head()
```

```
Out[45]:
```

	Date	TV	Radio	Newspaper	Sales
0	Thursday, January 1, 2015	230.1	37.8	69.2	22.1
1	Thursday, January 8, 2015	44.5	39.3	45.1	10.4
2	Thursday, January 15, 2015	17.2	45.9	69.3	9.3
3	Thursday, January 22, 2015	151.5	41.3	58.5	18.5
4	Thursday, January 29, 2015	180.8	10.8	58.4	12.9

```
In [8]: # Convert column index to list for readability
list(df.columns)
```

```
Out[8]: ['Date', 'TV', 'Radio', 'Newspaper', 'Sales']
```

```
In [9]: # Overview of dataset structure and data types
df.info()

# Statistical summary of numerical features
df.describe()

# Check for missing or null values
df.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 346 entries, 0 to 345
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Date        346 non-null    object  
 1   TV          346 non-null    float64 
 2   Radio        346 non-null    float64 
 3   Newspaper    346 non-null    float64 
 4   Sales        346 non-null    float64 
dtypes: float64(4), object(1)
memory usage: 13.6+ KB
```

```
Out[9]: Date      0
         TV       0
         Radio     0
         Newspaper 0
         Sales     0
         dtype: int64
```

### Formatting Issue

We can see that the date type is an object not an actual date format

Lets correct this first

```
In [10]: # Convert it to a proper datetime format for time-based analysis

df['Date'] = pd.to_datetime(df['Date'])

# Verify the conversion
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 346 entries, 0 to 345
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype    
--- 
 0   Date        346 non-null    datetime64[ns]
 1   TV          346 non-null    float64  
 2   Radio        346 non-null    float64  
 3   Newspaper    346 non-null    float64  
 4   Sales        346 non-null    float64 
dtypes: datetime64[ns](1), float64(4)
memory usage: 13.6 KB
```

```
In [11]: # Preview the converted datetime values

df['Date'].head()
```

```
Out[11]: 0    2015-01-01
         1    2015-01-08
         2    2015-01-15
         3    2015-01-22
         4    2015-01-29
Name: Date, dtype: datetime64[ns]
```

```
In [12]: df['Date'].isnull().sum() # ----- Currently we are checking the date format is correct
```

```
Out[12]: 0
```

## Identify the Constraint (CORE TOC STEP)

```
In [13]: # Compute correlation of advertising channels with Sales
# This helps identify the directional strength of each channel's impact

sales_corr = (
    df[['TV', 'Radio', 'Newspaper', 'Sales']]
    .corr()['Sales']
    .sort_values(ascending=False)
    * 100
)

sales_corr
```

```
Out[13]: Sales      100.000000
          Radio     38.410415
          TV        33.025625
          Newspaper  8.529707
          Name: Sales, dtype: float64
```

### Visual Validation of Directional Signal Using Regression Lines

- Now, lets check this out if we put some values/invest on any of this medium or channel
- which one is giving us most of the returns or generating the sales

```
In [14]: # Advertising Spend vs Sales
# Regression Visualization

fig, axes = plt.subplots(1, 3, figsize=(18, 5))

# TV vs Sales
sns.regplot(
    x='TV',
    y='Sales',
    data=df,
    ax=axes[0],
    scatter_kws={'alpha': 0.6},
    line_kws={'linewidth': 2}
)
axes[0].set_title('TV vs Sales\nStrong Linear Association')
axes[0].set_xlabel('TV Advertising Spend')
axes[0].set_ylabel('Sales')

# Radio vs Sales
sns.regplot(
    x='Radio',
```


```

        y='Sales',
        data=df,
        ax=axes[1],
        scatter_kws={'alpha': 0.6},
        line_kws={'linewidth': 2}
    )
axes[1].set_title('Radio vs Sales\nModerate Linear Association')
axes[1].set_xlabel('Radio Advertising Spend')
axes[1].set_ylabel('Sales')

# Newspaper vs Sales
sns.regplot(
    x='Newspaper',
    y='Sales',
    data=df,
    ax=axes[2],
    scatter_kws={'alpha': 0.6},
    line_kws={'linewidth': 2}
)
axes[2].set_title('Newspaper vs Sales\nWeak Linear Association')
axes[2].set_xlabel('Newspaper Advertising Spend')
axes[2].set_ylabel('Sales')

plt.tight_layout()
plt.show()

```



- As we can see, for TV & Radio there is an upward line which means if we are spending money on both they are giving us increase in sales compared to Newspaper

#### Takeaway

- While in the Newspaper the line is flat which means we are not able to generate much sales in the Newspaper.
- So our constraint is "Newspaper"

In [15]: # Multivariate Regression Analysis

```

# While correlation analysis and scatter plots provide directional insight,
# they do not quantify marginal throughput or isolate the individual impact
# of each advertising channel.
#
# A multivariate OLS regression is therefore performed to estimate the

```

```
# marginal contribution of each channel to Sales while controlling for
# the others.
#
# From a Theory of Constraints (TOC) perspective, this step is critical to
# distinguish true throughput drivers from channels that consume budget
# without meaningfully increasing output.

import statsmodels.api as sm

# Define predictors and target
X = df[['TV', 'Radio', 'Newspaper']]
X = sm.add_constant(X) # Add intercept term
y = df['Sales']

# Fit OLS model
model = sm.OLS(y, X).fit()

# Model summary
model.summary()
```

Out[15]:

OLS Regression Results

<b>Dep. Variable:</b>	Sales	<b>R-squared:</b>	0.260			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.253			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	39.97			
<b>Date:</b>	Tue, 06 Jan 2026	<b>Prob (F-statistic):</b>	3.64e-22			
<b>Time:</b>	01:41:02	<b>Log-Likelihood:</b>	-1072.2			
<b>No. Observations:</b>	346	<b>AIC:</b>	2152.			
<b>Df Residuals:</b>	342	<b>BIC:</b>	2168.			
<b>Df Model:</b>	3					
<b>Covariance Type:</b>	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	7.0526	0.823	8.566	0.000	5.433	8.672
<b>TV</b>	0.0244	0.003	7.194	0.000	0.018	0.031
<b>Radio</b>	0.1621	0.020	8.104	0.000	0.123	0.201
<b>Newspaper</b>	0.0009	0.011	0.081	0.935	-0.021	0.023
<b>Omnibus:</b>	9.896	<b>Durbin-Watson:</b>		2.118		
<b>Prob(Omnibus):</b>	0.007	<b>Jarque-Bera (JB):</b>		15.851		
<b>Skew:</b>	0.160		<b>Prob(JB):</b>	0.000361		
<b>Kurtosis:</b>	3.998		<b>Cond. No.</b>	497.		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## What can we understand from this?

- 
- 1. this is basically a Ordinary Least Squares multiple linear regression
  - 2. in Maths:  $\text{Sales} = \beta_0 + \beta_1(\text{TV}) + \beta_2(\text{Radio}) + \beta_3(\text{Newspaper}) + \epsilon$

In [16]: # Throughput Efficiency (Sales per Spend)

```
(df[['TV', 'Radio', 'Newspaper']]
 .div(df['Sales'], axis=0)
 .mean()
 .sort_values())
```

```
Out[16]: Radio      1.989754  
Newspaper    3.647249  
TV          12.372442  
dtype: float64
```

```
In [40]: # We have found the Constraint
```

## EXPLOIT THE CONSTRAINT

```
In [18]: df['Newspaper_bin'] = pd.qcut(df['Newspaper'], q=4)  
df.groupby('Newspaper_bin')['Sales'].mean()
```

```
Out[18]: Newspaper_bin  
(0.299, 18.3]      14.013636  
(18.3, 34.45]     13.672941  
(34.45, 57.45]    15.034884  
(57.45, 114.0]    15.736782  
Name: Sales, dtype: float64
```

```
In [19]: df['Newspaper'].quantile([0.25, 0.50, 0.75])
```

```
Out[19]: 0.25    18.30  
0.50    34.45  
0.75    57.45  
Name: Newspaper, dtype: float64
```

```
In [39]: df_exploit = df[df['Newspaper'] <= 35]
```

```
In [21]: df[['Sales']].mean(), df_exploit[['Sales']].mean()
```

```
Out[21]: (Sales    14.617052  
         dtype: float64,  
         Sales    13.820904  
         dtype: float64)
```

```
In [22]: df['Newspaper'].mean() - df_exploit['Newspaper'].mean()
```

```
Out[22]: 21.315646125208186
```

```
In [23]: df['Newspaper'].quantile([0.25, 0.50, 0.75])
```

```
Out[23]: 0.25    18.30  
0.50    34.45  
0.75    57.45  
Name: Newspaper, dtype: float64
```

```
In [24]: df_exploit = df[df['Newspaper'] <= 35]
```

```
In [25]: df['Sales'].mean(), df_exploit['Sales'].mean()
```

```
Out[25]: (14.617052023121387, 13.820903954802262)
```

```
In [26]: df['Newspaper'].mean() - df_exploit['Newspaper'].mean()
```

```
Out[26]: 21.315646125208186
```

```
In [38]: df['TP_ratio'] = df['Sales'] / df[['TV', 'Radio', 'Newspaper']].sum(axis=1)
df_exploit['TP_ratio'] = df_exploit['Sales'] / df_exploit[['TV', 'Radio', 'Newspaper']].sum(axis=1)
df['TP_ratio'].mean(), df_exploit['TP_ratio'].mean()
```

```
Out[38]: (0.08248236802063037, 0.09054665560726316)
```

```
In [28]: df_toc = df[df['Newspaper'] <= 35]
```

```
In [29]: df_toc[['TV', 'Radio', 'Sales']].corr()['Sales'].sort_values(ascending=False)
```

```
Out[29]: Sales      1.000000
          Radio     0.526301
          TV        0.433305
          Name: Sales, dtype: float64
```

```
In [30]: X = df_toc[['TV', 'Radio']]
X = sm.add_constant(X)
y = df_toc['Sales']
sm.OLS(y, X).fit().summary()
```

Out[30]:

## OLS Regression Results

<b>Dep. Variable:</b>	Sales	<b>R-squared:</b>	0.477			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.471			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	79.41			
<b>Date:</b>	Tue, 06 Jan 2026	<b>Prob (F-statistic):</b>	3.12e-25			
<b>Time:</b>	01:41:02	<b>Log-Likelihood:</b>	-492.39			
<b>No. Observations:</b>	177	<b>AIC:</b>	990.8			
<b>Df Residuals:</b>	174	<b>BIC:</b>	1000.			
<b>Df Model:</b>	2					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>const</b>	5.5141	0.735	7.499	0.000	4.063	6.965
<b>TV</b>	0.0287	0.004	8.163	0.000	0.022	0.036
<b>Radio</b>	0.2053	0.021	9.815	0.000	0.164	0.247
<b>Omnibus:</b>	18.380	<b>Durbin-Watson:</b>		2.175		
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>		61.450		
<b>Skew:</b>	0.230	<b>Prob(JB):</b>		4.53e-14		
<b>Kurtosis:</b>	5.850	<b>Cond. No.</b>		425.		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [31]:

```
df_toc['TV_Radio'] = df_toc['TV'] * df_toc['Radio']
X = df_toc[['TV', 'Radio', 'TV_Radio']]
X = sm.add_constant(X)
sm.OLS(y, X).fit().summary()
```

Out[31]:

## OLS Regression Results

<b>Dep. Variable:</b>	Sales	<b>R-squared:</b>	0.491			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.483			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	55.73			
<b>Date:</b>	Tue, 06 Jan 2026	<b>Prob (F-statistic):</b>	2.94e-25			
<b>Time:</b>	01:41:02	<b>Log-Likelihood:</b>	-489.94			
<b>No. Observations:</b>	177	<b>AIC:</b>	987.9			
<b>Df Residuals:</b>	173	<b>BIC:</b>	1001.			
<b>Df Model:</b>	3					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>const</b>	7.0080	0.995	7.047	0.000	5.045	8.971
<b>TV</b>	0.0183	0.006	3.136	0.002	0.007	0.030
<b>Radio</b>	0.1318	0.039	3.356	0.001	0.054	0.209
<b>TV_Radio</b>	0.0005	0.000	2.202	0.029	5.35e-05	0.001
	<b>Omnibus:</b>	28.856	<b>Durbin-Watson:</b>	2.210		
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	90.929			
<b>Skew:</b>	0.587	<b>Prob(JB):</b>	1.80e-20			
<b>Kurtosis:</b>	6.310	<b>Cond. No.</b>	1.39e+04			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.39e+04. This might indicate that there are strong multicollinearity or other numerical problems.

In [32]:

```
df_sim = df_exploit.copy()
df_sim['TV'] += 0.7 * 21.3
df_sim['Radio'] += 0.3 * 21.3


df_sim['Sales_pred'] = (
    7.008 +
    0.0183 * df_sim['TV'] +
    0.1318 * df_sim['Radio'] +
    0.0005 * df_sim['TV'] * df_sim['Radio']
)
df_sim['Sales_pred'].mean()
```

```
Out[32]: 15.552941232485875
```

```
In [33]: model.conf_int().loc['Newspaper']
```

```
Out[33]: 0   -0.021150  
1    0.022974  
Name: Newspaper, dtype: float64
```

```
In [34]: sm.qqplot(model.resid, line='45')  
plt.title("Residual Normality Check")  
plt.show()
```



```
In [41]: (df_sim['Sales_pred'].mean() - df_exploit['Sales'].mean()) / 21.3
```

```
Out[41]: 0.08131630411660155
```

## CONCLUSION

The analysis identifies Newspaper advertising as a low-impact constraint, while TV and Radio act as true sales drivers. Reallocation of budget from Newspaper to TV and Radio improves sales throughput without increasing total spend, supporting data-driven decision-making using the Theory of Constraints.